

Towards the Characterization of Realistic Model Generators using Graph Neural Networks

1st Given Name Surname
dept. name of organization (of Aff.)
name of organization (of Aff.)
City, Country
email address or ORCID

2nd Given Name Surname
dept. name of organization (of Aff.)
name of organization (of Aff.)
City, Country
email address or ORCID

Abstract—The automatic generation of software models is an important element in many software and systems engineering scenarios such as software tool certification, validation of cyber-physical systems, or benchmarking graph databases. Several model generators are nowadays available, but the topic of whether they generate realistic models has been little studied. The state-of-the-art approach to check the *realistic* property in software models is to rely on simple comparisons using graph metrics and statistics. This generates a bottleneck due to the compression of all the information contained in the model into a small set of metrics. Furthermore, there is a lack of interpretation in these approaches since there are no hints of why the generated models are not realistic. Therefore, in this paper, we tackle the problem of assessing how realistic a generator is by mapping it to a classification problem in which a Graph Neural Network (GNN) will be trained to distinguish between the two sets of models (real and synthetic ones). Then, to assess how realistic a generator is we perform the Classifier Two-Sample Test (C2ST). Our approach allows for interpretation of the results by inspecting the attention layer of the GNN. We use our approach to assess four state-of-the-art model generators applied to three different domains. The results show that none of the generators can be considered realistic.

Index Terms—Model generators, Realistic models, Graph neural networks, Two-Sample Test

I. INTRODUCTION

Model generators aim to create synthetic models automatically by taking one or more input parameters to configure the expected bounds or shape of the generated models. These types of tools can be applied to many areas of software and system engineering such as the testing and benchmarking of graph databases [1], [2], to create complex test stubs in the object-oriented field [1], [3] or automated synthesis of prototypical test contexts in the assurance of smart cyber-physical systems [1], [4], [5].

Recent works have established four properties that a model generator should satisfy [1], [5]: 1) *consistency* (the generator creates consistent models which satisfy all well-formedness constraints), 2) *diversity* (the generated models include a sufficiently wide variety of shapes [6]), 3) *scalable* (with respect to the size of a generated model) and 4) *realistic* (models generated cannot be distinguished from the real ones). In particular, a generator is *structurally realistic* if the set of generated models cannot be distinguished from the real ones just by looking at the typed graph structure (ignoring

the attribute values) [5], [7], [8]. Currently, a set of graph metrics and statistics (such as out-degree, dimensional degree, multiplex participation coefficient, etc) are used to assess whether a set of models can be considered similar to a dataset of realistic models [5], [7], [8]. This technique has three important shortcomings. Firstly, summarizing an entire graph model into a set of graph metrics causes an information loss. Secondly, a subset of graph statistics has to be chosen to perform the assessment, but not all metrics are equally effective to perform this task [7]. Thirdly, this approach is not interpretable, in the sense that it does not give us hints to determine why the generator is not realistic.

In this paper, we address the task of determining whether a model generator is realistic using a different technique, which overcomes these aforementioned issues. Our approach follows this idea: given two set of models, one generated by a given generator and the other composed by real models. Under the supposition that the generator is realistic (i.e., the synthetic models are realistic), if we build a Graph Neural Network (GNN) [9] trained to distinguish between these two sets, it will not be able to achieve a good performance since the classification problem is impossible. In order to evaluate how realistic a generator is, we use the non-parametric test called Classifier Two-sample Test (C2ST) [10]. The test tells us whether the GNN actually distinguished between real and generated models. Following this approach, we have assessed how realistic four state-of-the-art generators are in three domains. Altogether, this paper presents a novel approach for assessing model generators, which is more robust than previous approaches and has the additional advantage of being interpretable (i.e., by inspecting the weights of the attention layer of the GNN). Moreover, to the best of our knowledge, this is the first work that applies GNN to software models. Thus, our proposed GNN architecture can be adapted to face other model classification problems (e.g., meta-model classification [11], [12], UML classification [13], etc).

Organization. Section II gives a brief explanation of the technical background that underlies our approach, which is described in detail in Sect. III. In Sect. IV, we report our experiments assessing four state-of-the-art generators for three domains. Section V presents the main contributions of this

In this work, we refine the notion of realistic generator as follows: *Given a classifier trained with a sample of realistic and synthetic models, a generator is realistic if the classifier does not achieve a good performance distinguishing between synthetic and real models.*

C. Graph neural networks

A Graph Neural Network (GNN) or Graph Convolutional Neural Network (GCNN) [9] is a type of neural network that receives as input the nodes of a graph and generates node embeddings based on local network neighborhoods. More concretely, given a node $v \in V$ and its initial vector representation x_v , a GNN of L layers consist on the calculus of the final node embedding in this way:

$$h_v^0 = x_v \quad (1)$$

$$h_v^l = g_l(h_v^{l-1}, \{h_w^{l-1} : w \in \mathcal{N}(v)\}) \quad (2)$$

for $l = 1, \dots, L$, $\mathcal{N}(v) = \{w \in V | \exists e \text{ such that } f(e) = (w, v)\}$ (neighborhood of v) and g_l is a non-linear function that computes h_v^l using the vectors h_v^{l-1} and $\{h_w^{l-1} : w \in \mathcal{N}(v)\}$. Therefore, the final embedding h_v^L will have information of all neighbors in the L -hop neighborhood of v . To summarize, a GNN receives a set of nodes V in a graph and outputs node embedding vectors $\{h_v^L\}_{v \in V}$ that contain information of the L -hop neighborhood of each node.

The application of this type of layer can be seen as a message passing layer in which the embedding of a node is calculated by using the messages that it receives from its neighborhood. For example, Fig. 4 shows the application of a graph convolutional layer on our running example. Thick arrows are the messages that the vector v receives from its neighbors. Therefore, the new embedding associated to the node v (i.e., h_v^l) is calculated by applying g_l to its last embedding h_v^{l-1} and the embeddings of its neighborhood $h_{w_1}^{l-1}, h_{w_2}^{l-1}, h_{w_3}^{l-1}$ and $h_{w_4}^{l-1}$ (dotted arrows).

GNNs are used in problems such as node classification, link prediction or graph classification [22]. We are interested in the last one. In this task, the embedding vectors $\{h_v^L\}_{v \in V}$ are normally summarized in one vector:

$$h_G = \text{AGG} \{h_v^L : v \in V\} \quad (3)$$

where AGG is an order-invariant operation such as vector average, max operation, attention mechanisms, etc. This last vector represents the entire graph. Finally, h_G will be the input to a fully connected neural network which performs the classification using the sigmoid activation in its last layer (or softmax if the problem has more than two classes).

D. Classifier Two-Sample Test

The main aim of a two-sample test is to assess whether two samples $x_1, \dots, x_n \sim P$ and $y_1, \dots, y_n \sim Q$ come from

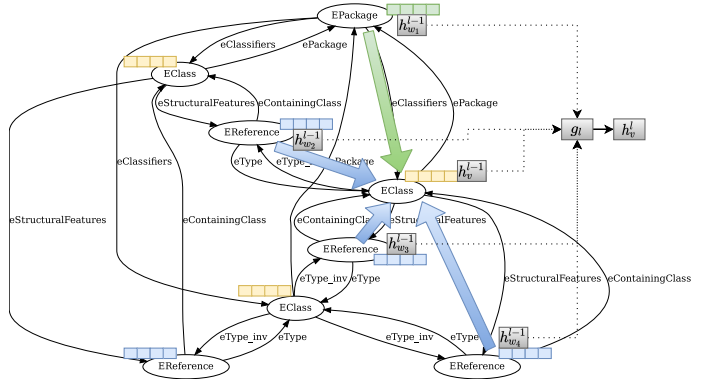


Fig. 4. Application of a layer in a GNN.

the same distribution, i.e., whether $P = Q$. In particular, the following test is considered:

$$\begin{aligned} H_0 : P &= Q, \\ H_1 : P &\neq Q \end{aligned}$$

The Classifier Two-Sample Test (C2ST) [10] tries to solve this problem using this idea: under H_0 (i.e., both distributions P and Q are the same), if we train a classifier to distinguish between P and Q (using the original samples), the classifier will not be able to achieve a good performance because both samples contain similar elements as they come from the same distribution. Thus, the expected accuracy (proportion of correctly predicted data samples) will be ~ 0.5 (near chance-level).

In details using the notation of [10], let us consider the dataset constructed using the samples $\{x_i\}_{i=1}^n$ and $\{y_i\}_{i=1}^n$ and associating a label to each sample (0 if it comes from P and 1 if it comes from Q):

$$\mathcal{D} = \{(x_i, 0)\}_{i=1}^n \cup \{(y_i, 1)\}_{i=1}^n = \{(z_i, l_i)\}_{i=1}^{2n}.$$

\mathcal{D} is randomly split into $\mathcal{D}_{\text{train}}$ and $\mathcal{D}_{\text{test}}$. After that, a classifier is trained using $\mathcal{D}_{\text{train}}$ to distinguish between labels 0 and 1. This model has the form $h(z) = P(l = 1|z)$ and it estimates the probability of a sample belonging to label 1. Then, it is evaluated on $\mathcal{D}_{\text{test}}$. The accuracy, \hat{t} , will be the statistic used to perform the hypothesis test:

$$\hat{t} = \frac{1}{|\mathcal{D}_{\text{test}}|} \sum_{(z_i, l_i) \in \mathcal{D}_{\text{test}}} \mathcal{I}(\mathcal{I}(h(z_i) > 0.5) = l_i), \quad (4)$$

where \mathcal{I} denotes the indicator function. Basically, \hat{t} is the proportion of correctly predicted data samples in the test set. The idea of C2ST is that, under H_0 , \hat{t} should be close to 0.5, but if we assume H_1 , \hat{t} should be greater than 0.5.

Following the procedure of the framework of statistical hypothesis testing [23], once the statistic \hat{t} has been computed, we are interested in the p -value, i.e., $P(T \geq \hat{t} | H_0)$. Under H_0 , the null distribution can be approximated by

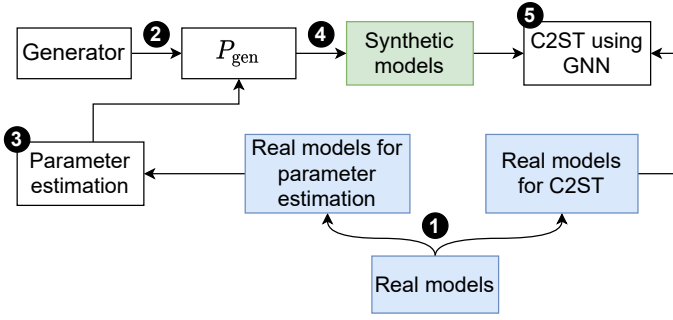


Fig. 5. Schema of the proposed approach.

$\mathcal{N}\left(\frac{1}{2}, \frac{1}{4|\mathcal{D}_{\text{test}}|}\right)$ [10]. Therefore, the p -value can be computed in practice using this formula:

$$P(T \geq \hat{t}|H_0) \approx \Phi\left(-\frac{\hat{t} - 0.5}{\sqrt{\frac{1}{4|\mathcal{D}_{\text{test}}|}}}\right). \quad (5)$$

Where Φ is the cumulative distribution function of a $\mathcal{N}(0, 1)$ distribution. If $p\text{-value} < \alpha$ (significance level), then H_0 ($P = Q$) is rejected, and accepted otherwise. Typically, α is set to 0.01 or 0.05.

III. APPROACH

Our approach to assess whether a generator is realistic or not is depicted in Fig. 5. The first step (label 1) is to split the dataset of real models in two sets (one used to perform the parameter estimation and the other one to perform the C2ST). Given a model generator, we make the assumption that the generated models are drawn (implicitly or explicitly) from some probability distribution to obtain diversity in the output models. Therefore, the second step in our approach is to map a model generator into a probability distribution over models (P_{gen} in Fig. 5 2). To achieve that, we have to estimate the parameters of the generator (which are reflected in its associated probability distribution) in order to force them to generate models that are as close as possible to the real ones (label 3). Using these parameters we generate a set of synthetic models (label 4). The assessment (label 5) is performed using CS2T by relying on a GNN-based classifier trained using a subset of real models and a set of synthetic models generated by the P_{gen} .

A. Model generators as distributions over models

Given a model generator, we need to determine which is its induced distribution (P_{gen}) from which models are sampled. This is a manual process which needs to be done for each particular generator by examining its features. In the following, we describe these processes for the four generators considered in this paper.

1) *EMF random instantiator* (RANDOM): This generator produces random instances for EMF meta-models [18]. It receives a meta-model \mathcal{M} as input, the number of objects (o) that the output must have and the average number of references per *EObject* (d). It does not support well-formedness constraints.




Fig. 6. Rules in RandomEMF.

Therefore, some of the generated output models could not be valid. We can see this generator as the following distribution over models (applying the law of total probability):

$$P_{\text{gen}}(M) = \sum_{(o,d) \in O} P(o,d)P_{\text{RANDOM}}(M|\mathcal{M}, o, d).$$

Therefore, to sample a model from this distribution we have to sample $(o, d) \sim P$ and then use the generator to sample M . In practice, the distribution P is approximated using pairs of (o, d) from real models (see next section).

2) *VIATRA*: This generator [1] receives a meta-model \mathcal{M} , the scope (normally the number of objects, o) of the output models and a set of well-formedness constraints Ψ . VIATRA maps the generation problem into a search problem of consistent models. It uses a back-end graph solver which makes this generator scalable. Similar to the EMF random instantiator, VIATRA can be seen as the following distribution:

$$P_{\text{gen}}(M) = \sum_{o \in O} P(o)P_{\text{VIATRA}}(M|\mathcal{M}, \Psi, o).$$

Therefore, to sample a model from this distribution we have to sample $o \sim P$ and then use VIATRA to sample M .

3) *Alloy*: Alloy Analyzer [19], [20] is a SAT-based model finder that can be used as generator of consistent models. It receives the same input as VIATRA and maps the generation problem into logic problem. Therefore, this generator can be seen as the following distribution:

$$P_{\text{gen}}(M) = \sum_{o \in O} P(o)P_{\text{Alloy}}(M|\mathcal{M}, \Psi, o).$$

When sampling from $P_{\text{Alloy}}(M|\mathcal{M}, \Psi, o)$, we add a random amount of extra true statements (as it is done in [8], [24]) to prevent the solver from running deterministically.

4) *RandomEMF* (rEMF): This generator is a rule-based generator in which the generation process is driven by user-defined rules. It belongs to the category of generators that are based on formal and graph grammars [21]. There are two types of *random* rules, which are illustrated in Fig. 6. In this example, the first rule (called *Package*) corresponds to the root rule, and it indicates that an object of type *EPackage* must be generated and the number of its children classifiers must follow a negative binomial distribution with parameters 1 and 0.9784. The second rule is an *alternative* rule, and it can randomly derive in one of these three rules: *Enum*, *DataType* or *Class* with priorities 1, 4 and 22 respectively.

Since this generator receives a meta-model and a set of grammar rules (\mathcal{G}), it can be seen as the following distribution over models:

$$P_{\text{gen}}(M) = P_{\text{rEMF}}(M|\mathcal{M}, \mathcal{G}).$$

Similar to EMF random instantiator, this generator does not support well-formedness constraints. The set of grammar rules must be designed to enforce the generation of consistent models as much as possible. In our experiments, when designing the rules we have attempted to respect the majority of constraints. However, due to the limitations of the RandomEMF grammar definition language, we could not assure all constraints to be respected in the generated output models. For instance, in Ecore, there is no guarantee that the generated models will respect the restriction of *no cycles in hierarchy*.

B. Parameter estimation

To sample models from P_{gen} we need to estimate or approximate its parameters from a set of realistic models. This process is different depending on the parameters of the generators.

1) *Parameter estimation for VIATRA and Alloy*: To approximate $P(o)$ we use the Kernel Density Estimation (KDE), which is a well-known method to estimate the density function of a distribution. Let us suppose that we have a set of samples $\{o_1, \dots, o_n\}$ corresponding to the number of objects in each realistic model. As estimation of the density function KDE takes:

$$\hat{f}_{h,K}(o) = \frac{1}{nh} \sum_{i=1}^n K\left(\frac{o - o_i}{h}\right).$$

Where K is a kernel function and $h > 0$ is called bandwidth. K and h are hyperparameters and they are chosen using cross-validation over the set of samples. Once these hyperparameters are fixed, we can sample from $\hat{f}_{h,K}$ and obtain $\{o'_1, \dots, o'_m\} \sim \hat{f}_{h,K}$. Due to the fact that the new samples are not integers, we apply the floor function to them obtaining $\{\lfloor o'_1 \rfloor, \dots, \lfloor o'_m \rfloor\}$.

2) *Parameter estimation for RANDOM*: In this case, we need to approximate $P(o, d)$ since RANDOM depends on the number of objects and the average number of references. We apply the same idea as before, but to pairs (o, d) by using $\hat{f}_{h,K}$ in 2 dimensions.

3) *Parameter estimation in rEMF*: As explained before, to generate models with rEMF, we need to define a set of rules \mathcal{G} . Our approach is to manually define these rules for the relevant meta-model elements, and then we estimate the parameters of each rule individually. Since rEMF supports two types of rules, we need to estimate two types of parameters:

Shape of distribution. For normal rEMF rules, we are interested in determining which is the distribution that best represents the characteristics of each type of meta-model element in the set of realistic models. To do so, we estimate the parameters of each distribution associated to the rules. For each different rule, we build a set of samples. For instance,

for rule 1 in Fig. 6, a sample is the count of the number of classifiers of a *EPackage*. As a result, the set of samples $\{c_1, \dots, c_s\}$ is obtained. Then, for each available distribution in rEMF (Uniform, Normal, Negative Binomial and Poisson), we estimate its parameters using maximum likelihood. Finally we take the best distribution that fits $\{c_1, \dots, c_s\}$, using the log-likelihood as the comparison criteria.

Priority in alternative rules. In rule 2 of Fig. 6, the priorities of each one of the alternative rules (*Enum*, *Data Type*, *Class*) need to be fixed. To do so, for each model we count the proportion of classifiers that it has (i.e. *EEnum*, *EData Type* and *EClass*). Then, we compute the mean with respect to all models. Finally, these averages are divided by the minimum average. Since priorities are integers, we round them.

C. Applying GNNs to graph models

In the previous section we have explained how we estimate the parameters needed to force the generator to generate synthetic models that are as similar as possible to the dataset of real models. At this point, we have two sets of models (realistic models and synthetic models). We need a classifier to distinguish between these two sets in order to perform the C2ST. Therefore, we choose a GNN as our classifier. This neural network will receive a model (previously converted into a graph as described in Sect. II-A) as input and determine if it is real or not. Since we are dealing with graphs with labeled edges and nodes, we will consider the PyTorch [25], [26] implementation of the graph convolutional layer proposed in [27]. Therefore, the node embeddings are calculated using the following formulas:

$$h_v^0 = E_{\text{node}}(v) \quad (6)$$

$$h_v^l = \text{ReLU} \left(W_l^1 h_v^{l-1} + \sum_{r \in \mathcal{R}} \sum_{w \in \mathcal{N}_r(v)} \frac{1}{|\mathcal{N}_r(v)|} W_l^r h_w^{l-1} \right) \quad (7)$$

where E_{node} is the embedding layer of node labels (initial embedding of the nodes), \mathcal{R} is the set of references (edge types) and $\mathcal{N}_r(v)$ is the neighborhood of v restricted to r . Note that equations 6 and 7 are adaptations of equations 1 and 2 respectively. In particular, the initial node embedding is given by an embedding layer and g_l is a function that takes into account the labels of the edges.

In order to have an embedding at the graph level, we will use an attention vector to summarize all the node embeddings i.e., the AGG operator in Eq. 3 is defined as follows:

$$h_G = \sum_{v \in V} \alpha_v h_v^L \text{ where } \alpha_v = \frac{\exp(\alpha \cdot h_v^L)}{\sum_{w \in V} \exp(\alpha \cdot h_w^L)}. \quad (8)$$

Finally, since we are interested in a binary classification problem, the output of the model is computed as

$$P_{\text{model}}(\text{is real?}|G) = \sigma(W_F \cdot h_G + b_F),$$

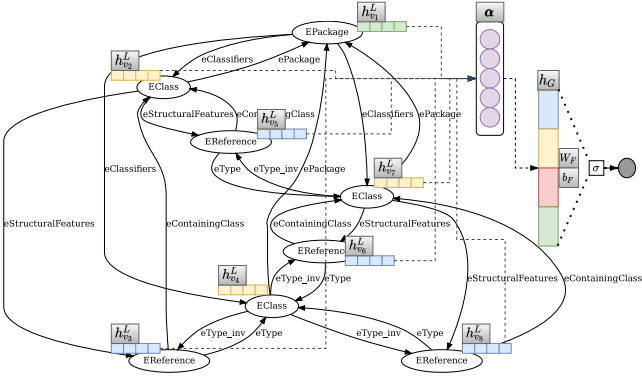


Fig. 7. Attention mechanism and last linear layer.




Fig. 8. Training phase and global architecture of the neural network.

where σ is the sigmoid function. Fig. 7 shows the attention mechanism and the last linear layer applied to the running example. There, once the graph convolutional layers are applied, each node v_1, \dots, v_8 has an embedding associated $h_{v_1}^L, \dots, h_{v_8}^L$. Each $h_{v_i}^L$ contains the information of the L -hoop neighborhood of v_i . To obtain the embedding of the entire graph, a weighted averaged is performed over the node embeddings using Eq. 8 (dotted lines that connect each $h_{v_i}^L$ with α). The greater the weight associated to a node is, the more influence it has on the output. For example, if the attention weight associated to the node v_2 is 0.3 while the weights associated to the rest are 0.1, then the neural network focuses more in v_2 and its L -hoop neighborhood than in other nodes. Finally, h_G is passed through a dense layer followed (W_F and b_F) by a sigmoid function (σ) to output probabilities.

In our experiments, all dimensions of the layers and vectors are fixed to 64. We take $L = 2$ since the diameter of the inputs graphs is relatively small. With this setting, the GNN takes into account the 2-hoop neighborhood. Therefore, all metrics studied in [7] are implicitly considered by our neural model. Finally, the neural network is trained using the binary cross-entropy loss and the Adam optimizer [28].

The training phase of the final model is depicted in Fig. 8. Each node of the graph is mapped to an initial vector by using the embedding layer E_{node} . These vectors are passed through a two layer GNN. After that, they are summarized into a vector that represents the entire graph by using the attention vector α . The graph vector is the input of a fully connected layer (W_F, b_F) that outputs probabilities. Finally, the output is compared with the ground truth (using the binary cross-

entropy loss) and all weights of the neural model are updated (i.e., weights of E_{node} , GNN_1 , GNN_2 , α , W_F and b_F) using backpropagation and the Adam optimizer.

D. Assessing realistic generators

Putting all together, a generator (gen) is a black box that receives as input a set of conditions (that can be a meta-model, some restrictions, etc) and outputs or samples a model M (that is consistent with the input conditions, i.e., conforms to a meta-model, satisfy some restrictions, etc). As it was explained, this black box can be seen as a probability distribution over models $P_{gen}(M)$. We can suppose that all realistic models (models that people make) are sampled from a probability distribution Q_{real} . Thus, we can say that a generator is realistic if P_{gen} is similar to Q_{real} .

In practice, we have a set of real models $\{R_1, \dots, R_n\}$ sampled from Q_{real} and a set of models $\{S_1, \dots, S_n\}$ sampled from P_{gen} . To determine if $P_{gen} \approx Q_{real}$, we apply the following hypothesis test:

$$H_0 : P_{gen} = Q_{real}$$

$$H_1 : P_{gen} \neq Q_{real}$$

Using C2ST (Sect. II-D) and a classifier (the GNN model explained in Sect. III-C), we can compute the accuracy statistic and the p -value. If p -value $< \alpha$, then we reject H_0 and the generator is not realistic. The bigger p -value the better. Moreover, if the accuracy is ~ 0.5 , the generator outputs realistic models. On the other hand, if the accuracy is high, it means that the generator is not realistic with respect to the dataset.

The application of this procedure to assess concrete model generators is explained in detail in the next section.

IV. EXPERIMENTS

In this section we report the results of our experiments applying our approach to assess whether state-of-the-art model generators are able to produce realistic models. We focus on the generators already explained in Sect. III-A, that is, EMF random instantiator [18], RandomEMF [21], VIATRA [1] and Alloy [19], [20].

A. Domains

Our experiments were based on three case studies, each one corresponding to a different domain. The selection of the case studies was driven by the availability of public models. In particular, we obtained these models from the MAR search engine¹ [16].

- **Ecore.** We considered a reduced version of Ecore [14], as depicted in Fig. 1, which included six well-formedness constraints (denoted Ψ in Sect. III-A). Then, we downloaded 500 real models from MAR (which were crawled from GitHub) and transformed them into our reduced version. To ensure that the generators could generate similar models, we chose models that verify the following

¹<http://mar-search.org>

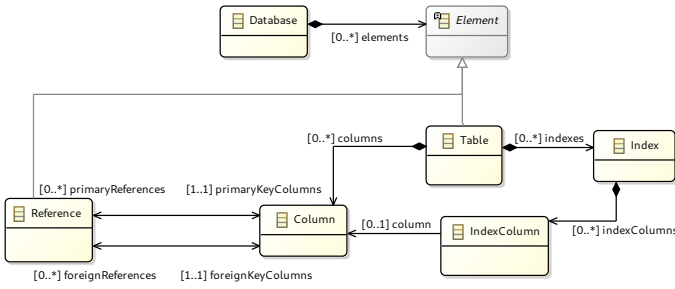


Fig. 9. Simplified version of the database model meta-model in GenMyModel.

two conditions: the model must have just one *EPackage* as root and the model does not contain attributes of type *EInt*, *EString*, etc.

- **Yakindu Statecharts** [29] is an industrial modeling environment. We reused the meta-model and the ten well-formedness constraints defined in [5]. Regarding the dataset real models, a set of 369 has been considered (which were crawled from GitHub).
- **Database models** from GenMyModel². We considered a manually constructed version of its meta-model (depicted Fig. 9) that includes tables, indices and references. Regarding well-formedness constraints, we manually derived three constraints according to our observations using the platform. Then, we built a dataset of 500 real models.

B. Procedure

Given an initial set of real models \mathcal{R} and a generator, each experiment is performed by following these steps, which are illustrated in Fig. 10:

- 1) \mathcal{R} is split into \mathcal{R}_I (50%) and \mathcal{R}_{II} (50%). This step corresponds to label ① in Fig. 10. The goal is to ensure that the assessment procedure (C2ST) is not biased by the parameter estimation.
- 2) \mathcal{R}_{II} is used to estimate some parameters of the generator (i.e., $P(o)$ in VIATRA and Alloy, $P(o, d)$ in RANDOM, shapes and priorities of \mathcal{G} in rEMF). This step corresponds to label ② in Fig. 10.
- 3) Once the parameters of the generator are fixed, the generator is used to generate $n = |\mathcal{R}_I|$ models, which form the set of synthetic models \mathcal{S} (label ③ in Fig. 10).
- 4) The sets \mathcal{R}_I and \mathcal{S} are merged and shuffled (label ④).
- 5) Finally, we apply C2ST considering that $\mathcal{S} \sim P_{\text{gen}}$ and $\mathcal{R}_I \sim Q_{\text{real}}$ (label ⑤). To do so, the merged set of models is split into training/validation/test (60%/15%/25%), label ⑤.1. The training set is used to train the GNN model and the validation set is used to perform early stopping³ (label ⑤.2). As a result, a trained GNN model is obtained. This model is evaluated using the test

²<https://www.genmymodel.com/>

³Regularization mechanism that consists on stopping the training when there is no improvement over the validation set.

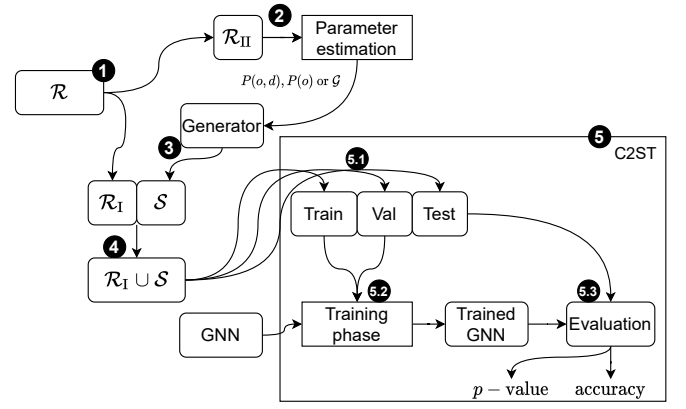


Fig. 10. Experimental procedure.

set (label ⑤.3). Finally, the accuracy and p -value are calculated using equations 4 and 5 respectively. The results are used to determine whether the generator is realistic or not.

C. Results

The results of the C2ST for each generator and for each domain are shown in Table I. According to them, we can draw the following conclusions:

- No generator can be considered realistic since all p -values are less than $\alpha = 0.01$. Therefore, we cannot accept $H_0 : P_{\text{gen}} = Q_{\text{real}}$.
- The least realistic generator is RANDOM. This is because it actually generates models that conform to a meta-model, but the majority of the generated models are inconsistent with the well-formedness constraints.
- The most realistic generator is rEMF. This could be justified for two reasons. First, rEMF is the most customizable generator since the user has to define the set of rules \mathcal{G} that guides the model generation. And second, there is more observation of the set \mathcal{R}_{II} since we estimate the parameters of many distributions (for example, in the Ecore domain, number of classifiers per *EPackage*, number of features per *EClass*, etc) to construct \mathcal{G} . However, in RANDOM, Alloy and VIATRA only $P(o, d)$ and $P(o)$ are set.

Altogether, these experiments show that the distributions of realistic models have a complexity which cannot be approximated by the distributions implemented by current model generators. Therefore, new techniques and procedures are needed to enable the faithful generation of realistic models.

D. Interpretation

The described procedure tells us whether a generator is realistic or not. However, faced with a negative answer we need some mechanism to interpret this result and understand it. Since our procedure is based on C2ST, we can determine which models of the test set have been correctly or incorrectly classified by the GNN. Then, for each interesting model, we

TABLE I
RESULTS OF C2ST.

DOMAIN	RANDOM		ALLOY		VIATRA		REMF	
	accuracy	p-value	accuracy	p-value	accuracy	p-value	accuracy	p-value
Ecore ($ \mathcal{D}_{\text{test}} = 125$)	1	2.544E-29	0.9600	4.078E-25	0.9280	5.326E-22	0.8880	2.050E-18
Yakindu ($ \mathcal{D}_{\text{test}} = 93$)	1	2.614E-22	0.9462	3.759E-18	0.9569	6.033E-19	0.7526	5.477E-07
Database ($ \mathcal{D}_{\text{test}} = 125$)	0.9920	1.880E-28	0.9200	2.957E-21	0.9680	6.267E-26	0.6880	1.312E-05

can inspect the behavior of the neural network when this particular model is introduced as input. Our GNN model has an attention mechanism that can be easily interpreted by just looking at its weights. This feature is particularly useful when the input model is large. In this case, we can focus only on the parts of the model that the GNN has focused on.

As a concrete example, let us consider the worst and best generators: EMF random instantiator and RandomEMF. For each generator we show one example model per domain which is not considered realistic by the corresponding GNN. We use a custom notation in which each object in the model is given a color according to the attention that the GNN has put on it (see Fig. 11 and Fig. 12). These graphics are automatically generated using specific tooling that we have developed.

EMF random instantiator is not considered realistic because its generated models violate the well-formedness constraints. In almost all the analyzed models, the GNN focuses on some object that violate some constraint. For instance, Fig. 11 shows three examples of invalid models. In the Ecore model (label **a**), the GNN focuses on a *EClass* that inherits from itself. In the Yakindu domain (label **b**), the model focuses on transitions whose target state is an entry (this violates a constraint). Finally, label **c** shows a synthetic database model, for which the GNN focuses on references whose target and source columns are the same column (this again violates a constraint).

Regarding RandomEMF, the criteria followed by the GNN to distinguish between generated and real model is more diverse (Fig. 12). For instance, in the case of Ecore (label **a**), the GNN focuses on *EEnum* objects that are not referenced by any *EAttribute* belonging to an *EClass* of the same package. The GNN correctly detects that it is an uncommon situation to create an *EEnum* but not using it in the meta-model. In Yakindu models (label **b**), the GNN identifies states that are isolated (i.e., they are not pointed to any transition). This is a common pattern in models generated by RandomEMF but not in real models. In database models (label **c**), we can observe that there are cases in which RandomEMF generates two objects *Reference* that are opposites i.e., the target column of one of them is the source column and vice versa. This pattern is allowed (you can instantiate it by using the editor and it does not break any constraint) but it does not make sense and it is not common in real models.

Data Availability. The replication package together with all attention heatmaps of the test models for each generator and for each domain are available at <https://drive.google.com/drive/folders/>

1Op2LRCdvhxsGrpTRoLT2AAZXXz96-z-f?usp=sharing.

V. ASSESSMENT

Our experiments show that our method is able to effectively address the task to characterizing a model generator in terms of the *realistic* property. Nevertheless, we found some limitations.

The main limitation of our approach is that it depends on the samples of real models. In the experiments, we assumed that the collected models are a representative sample of real models. However, they are only a sample of available public models. This was a threat to validity in our experiments, and we cannot claim that the evaluated generators are not generally able to produce realistic models. In practice, our approach will always be biased by the selection of the set of real models.

Our approach is based on C2ST. Therefore, it depends on the selected classifier. A very simple classifier could not be able to distinguish between two distributions even if they are the same. To mitigate this, we use GNN as classifier. This neural model has been successfully used in a variety of complicated graph classification problems [22].

There are cases in which the interpretation is difficult to perform just by inspecting the models in the test set and the attention weights. Moreover, the interpretation part of our approach is manual, so it can be time-consuming as one needs to figure out why the GNN has focused on certain nodes. As future work we plan to incorporate more sophisticated interpretation techniques, such as [30] or [31], that can make the interpretation easier and faster.

Despite these limitations, this work brings a new perspective to the field of model generators. The presented technique provides a robust and interpretable method to assess whether model generators are realistic. Moreover, the key elements of the approach are also useful contributions on themselves which can be useful for other purposes, namely:

- A method to interpret a set of models as samples from a probability distribution over models, combined with the use C2ST and GNNs to compare realistic and synthetic models. This can also be applied to other comparison problems, such as comparing two versions of different generators and comparing datasets of model (e.g., coming from different repositories).
- We have also devised a technique to determine which parameters of a given model generator are more likely to produce the closest models with respect to a given model dataset. This is useful to enable the automated tuning of a model generator.
- Moreover, to the best of our knowledge, this is the first work that applies GNN to software models. Thus,

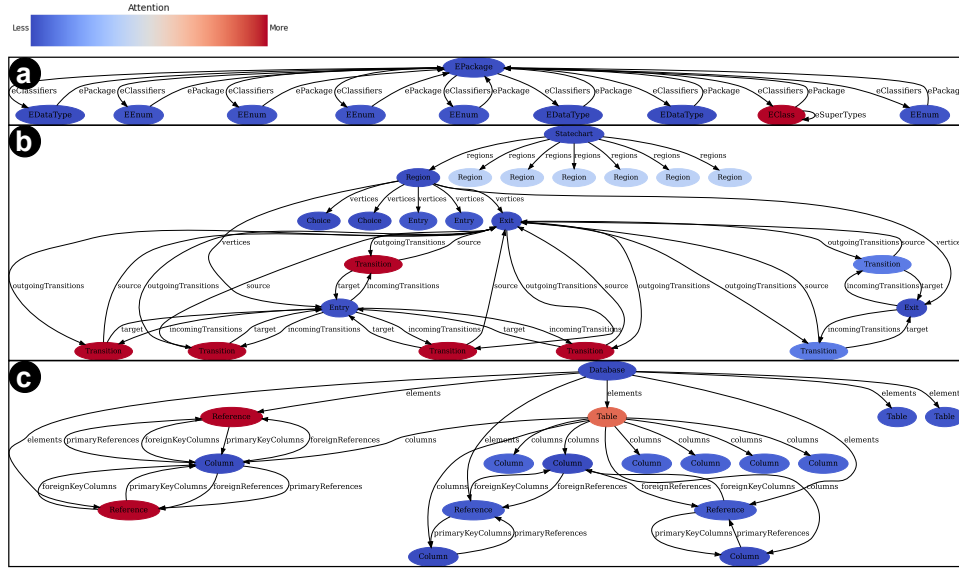


Fig. 11. Attention heatmap of models generated with EMF random instantiator

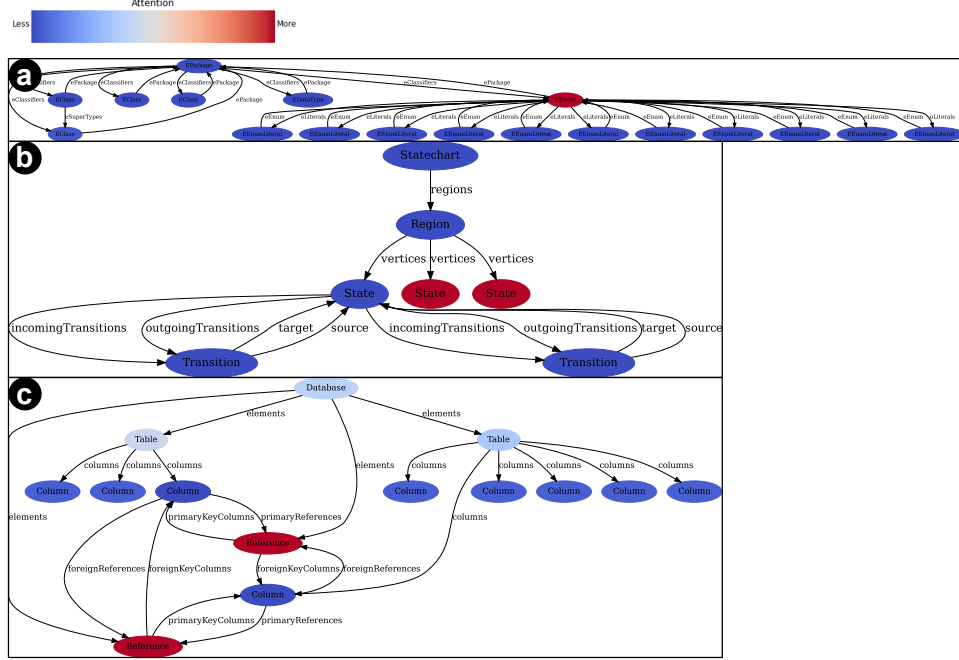


Fig. 12. Attention heatmap of models generated with RandomEMF

our proposed GNN architecture can be adapted to face other model classification problems (e.g., meta-model classification [11], [12], UML classification [13], etc).

This paper presents a novel procedure to assess realistic model generators. The approach is interpretable since it is based on the transparent test C2ST and uses a GNN with an attention mechanism.

VI. RELATED WORK

In this section we review works related to our proposal, organized in three categories: two-sample tests, assessment of

realistic generators and model generators.

A. Two-sample tests

The aim of these tests is to determine if two sets of samples come from the same distribution. Traditional examples are Student t -test and Kolmogorov-Smirnov (KS) test. Although these tests are widely used, they require strong assumptions about the data (e.g., normality assumption) [32]. Furthermore, they cannot be applied to discrete data (such as graphs or strings). In this line, the Kernel Two-Sample Test [33] is a powerful test based on kernels and the Maximum Mean

Discrepancy (MMD). It can be applied to any type of data as long as you can define a kernel function. On the other hand, another popular test is the Classifier Two-Sample Test [10]. The idea is that under the assumption of the null hypothesis, the classification problem of distinguishing between samples from a distribution and the other is impossible. It can be proved that this test is a particular case of the Kernel Two-Sample Test [32]. We use the Classifier Two-Sample Test in our approach.

B. Assessment of realistic model generators

In [7], the authors propose the use of multidisciplinary graph metrics to characterize realistic models. The work in [5] uses these metrics to compare two generators. Given a graph statistic (degree distribution, multiplex participation coefficient, pairwise multiplexity, etc), they measure how realistic a generator is by using the average value of the KS statistic between each pair of models (A, B) where A belongs to the set of real models and B to the set of generated models. This technique has three important shortcomings. Firstly, summarizing an entire graph into a set of graph metrics causes an information loss. Secondly, a subset of graph statistics have to be chosen to perform the assessment, but not all metrics are equally effective to perform this task [7]. Thirdly, this approach is not interpretable. Moreover, the KS based measure is not consistent since it does not take into account the differences inside the set of real models and inside the set of synthetic models. On the other hand, [8] claims that a generated model is realistic if its distance (considering graph metrics and statistics) with respect to its nearest neighbor real model is less than a threshold. To measure how realistic a generator is, the authors propose to take the average of all the distances and the percentage of realistic synthetic models. This proposed method is not consistent since it can be *hacked* by a generator that always generates the same realistic model. In our approach, the GNN learns the structure of realistic models by looking for differences between both synthetic and realistic sets of models, so it is robust to this scenario. An evaluation metric based on the Maximum Mean Discrepancy [33] is proposed in [34]. This metric is used to assess generative models of graphs and it is similar to the one proposed in [5] (except that it takes into account the differences inside the groups which makes it a consistent measure).

C. Model generators

Some model generators are based on mapping the meta-model and constraints to logical formulas and use SAT-solvers to obtain consistent models. It is possible to build this type of generators using model finders like Formula Framework [35], Alloy [20] and EMF2CSP [36]. In rule-based generators, the generation process is guided by rules. Examples of this type of generators are RandomEMF [21], [37] or [38]. Search-based generators transform the generation into a search problem. Examples are VIATRA [1] and [39]. Random models of graphs (such as Erdős-Rényi [40] or Watts-Strogatz [41]) can be considered model generators. However, they are mostly for

unlabeled graphs. Since there exist networks whose complexity is beyond these random models, some generative models of graphs have been proposed. For instance, GraphRNN [34] uses a two-level recurrent neural network to generate a graph. GraphVAE [42] uses variational autoencoders. [43] transforms the generation of a graph into a sequence of actions and uses a GNN to decide which action should be taken in each step. These generative models need a training set and they are able to generate graphs whose properties are close to the training set. They can be used in model generation because they can be adapted to generate labeled graphs.

VII. CONCLUSION

In this work we have presented a novel method to assess whether a given generator is realistic or not. Our approach is based on training a GNN in the task of distinguishing between real models and synthetic models. If the classifier does not achieve a good performance, the generator is realistic. The proposed method is interpretable and some hints of why the generator is not realistic can be derived by looking at the samples in the test set and the attention weights of the GNN model. On the other hand, the results of applying the proposed assessment to four state-of-the-art generators in three different domains show that none of the generators is realistic. Therefore, more research is needed to obtain realistic model generators.

As future work, we will try to automate the interpretation of our approach. Furthermore, we plan to include this approach in a generator in order to guide the generation process and obtain realistic models.

REFERENCES

- [1] O. Semeráth, A. S. Nagy, and D. Varró, "A graph solver for the automated generation of consistent domain-specific models," in *Proceedings of the 40th International Conference on Software Engineering*, 2018, pp. 969–980.
- [2] G. Bagan, A. Bonifati, R. Ciucanu, G. H. Fletcher, A. Lemay, and N. Advokaat, "gmark: Schema-driven generation of graphs and queries," *IEEE Transactions on Knowledge and Data Engineering*, vol. 29, no. 4, pp. 856–869, 2016.
- [3] D. Marinov and S. Khurshid, "Testera: A novel framework for automated testing of java programs," in *Proceedings 16th Annual International Conference on Automated Software Engineering (ASE 2001)*. IEEE, 2001, pp. 22–31.
- [4] Z. Micskei, Z. Szatmári, J. Oláh, and I. Majzik, "A concept for testing robustness and safety of the context-aware behaviour of autonomous systems," in *KES International Symposium on Agent and Multi-Agent Systems: Technologies and Applications*. Springer, 2012, pp. 504–513.
- [5] D. Varró, O. Semeráth, G. Szárnyas, and Á. Horváth, "Towards the automated generation of consistent, diverse, scalable and realistic graph models," in *Graph Transformation, Specifications, and Nets*. Springer, 2018, pp. 285–312.
- [6] O. Semeráth, R. Farkas, G. Bergmann, and D. Varró, "Diversity of graph models and graph generators in mutation testing," *International Journal on Software Tools for Technology Transfer*, vol. 22, no. 1, pp. 57–78, 2020.
- [7] G. Szárnyas, Z. Kővári, Á. Salánki, and D. Varró, "Towards the characterization of realistic models: evaluation of multidisciplinary graph metrics," in *Proceedings of the ACM/IEEE 19th International Conference on Model Driven Engineering Languages and Systems*, 2016, pp. 87–94.
- [8] O. Semeráth, A. A. Babikian, B. Chen, C. Li, K. Marussy, G. Szárnyas, and D. Varró, "Automated generation of consistent, diverse and structurally realistic graph models," *Softw. Syst. Model.*, 2021, in press.

- [9] F. Scarselli, M. Gori, A. C. Tsoi, M. Hagenbuchner, and G. Monfardini, "The graph neural network model," *IEEE transactions on neural networks*, vol. 20, no. 1, pp. 61–80, 2008.
- [10] D. Lopez-Paz and M. Oquab, "Revisiting classifier two-sample tests," *arXiv preprint arXiv:1610.06545*, 2016.
- [11] P. T. Nguyen, J. Di Rocco, D. Di Ruscio, A. Pierantonio, and L. Iovino, "Automated classification of metamodel repositories: A machine learning approach," in *2019 ACM/IEEE 22nd International Conference on Model Driven Engineering Languages and Systems (MODELS)*. IEEE, 2019, pp. 272–282.
- [12] P. T. Nguyen, D. Di Ruscio, A. Pierantonio, J. Di Rocco, and L. Iovino, "Convolutional neural networks for enhanced classification mechanisms of metamodels," *Journal of Systems and Software*, vol. 172, p. 110860, 2021.
- [13] M. H. Osman, T. Ho-Quang, and M. Chaudron, "An automated approach for classifying reverse-engineered and forward-engineered uml class diagrams," in *2018 44th Euromicro Conference on Software Engineering and Advanced Applications (SEAA)*. IEEE, 2018, pp. 396–399.
- [14] D. Steinberg, F. Budinsky, E. Merks, and M. Paternostro, *EMF: eclipse modeling framework*. Pearson Education, 2008.
- [15] R. Clarisó and J. Cabot, "Applying graph kernels to model-driven engineering problems," in *Proceedings of the 1st International Workshop on Machine Learning and Software Engineering in Symbiosis*, 2018, pp. 1–5.
- [16] J. A. H. López and J. S. Cuadrado, "Mar: A structure-based search engine for models," in *Proceedings of the 23rd ACM/IEEE International Conference on Model Driven Engineering Languages and Systems*, 2020, pp. 57–67.
- [17] Ö. Babur, L. Cleophas, and M. van den Brand, "Metamodel clone detection with samos," *Journal of Computer Languages*, vol. 51, pp. 57–74, 2019.
- [18] "Atlantmod team (inria, mines-nantes, lina), emf random instantiator." [Online]. Available: <https://github.com/atlanmod/mondo-atlzo-benchmark/tree/master/fr.inria.atlanmod.instantiator>
- [19] E. Torlak and D. Jackson, "Kodkod: A relational model finder," in *International Conference on Tools and Algorithms for the Construction and Analysis of Systems*. Springer, 2007, pp. 632–647.
- [20] D. Jackson, "Alloy: a lightweight object modelling notation," *ACM Transactions on Software Engineering and Methodology (TOSEM)*, vol. 11, no. 2, pp. 256–290, 2002.
- [21] M. Scheidgen, "Generation of large random models for benchmarking," *BigMDE@ STAF*, vol. 1406, pp. 1–10, 2015.
- [22] Z. Wu, S. Pan, F. Chen, G. Long, C. Zhang, and S. Y. Philip, "A comprehensive survey on graph neural networks," *IEEE transactions on neural networks and learning systems*, 2020.
- [23] G. Casella and R. L. Berger, *Statistical inference*. Cengage Learning, 2021.
- [24] O. Semeráth and D. Varró, "Iterative generation of diverse models for testing specifications of dsl tools," in *FASE*, vol. 18, 2018, pp. 227–245.
- [25] M. Fey and J. E. Lenssen, "Fast graph representation learning with pytorch geometric," *arXiv preprint arXiv:1903.02428*, 2019.
- [26] A. Paszke, S. Gross, F. Massa, A. Lerer, J. Bradbury, G. Chanan, T. Killeen, Z. Lin, N. Gimelshein, L. Antiga, A. Desmaison, A. Kopf, E. Yang, Z. DeVito, M. Raison, A. Tejani, S. Chilamkurthy, B. Steiner, L. Fang, J. Bai, and S. Chintala, "Pytorch: An imperative style, high-performance deep learning library," in *Advances in Neural Information Processing Systems* 32, H. Wallach, H. Larochelle, A. Beygelzimer, F. d'Alché-Buc, E. Fox, and R. Garnett, Eds. Curran Associates, Inc., 2019, pp. 8024–8035. [Online]. Available: <http://papers.neurips.cc/paper/9015-pytorch-an-imperative-style-high-performance-deep-learning-library.pdf>
- [27] M. Schlichtkrull, T. N. Kipf, P. Bloem, R. Van Den Berg, I. Titov, and M. Welling, "Modeling relational data with graph convolutional networks," in *European semantic web conference*. Springer, 2018, pp. 593–607.
- [28] D. P. Kingma and J. Ba, "Adam: A method for stochastic optimization," *arXiv preprint arXiv:1412.6980*, 2014.
- [29] "Yakindu statechart tools." [Online]. Available: <https://www.itemis.com/en/yakindu/state-machine/>
- [30] H. Yuan, J. Tang, X. Hu, and S. Ji, "Xggn: Towards model-level explanations of graph neural networks," in *Proceedings of the 26th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, 2020, pp. 430–438.
- [31] R. Ying, D. Bourgeois, J. You, M. Zitnik, and J. Leskovec, "Gnnexplainer: Generating explanations for graph neural networks," *Advances in neural information processing systems*, vol. 32, p. 9240, 2019.
- [32] F. Liu, W. Xu, J. Lu, G. Zhang, A. Gretton, and D. J. Sutherland, "Learning deep kernels for non-parametric two-sample tests," in *International Conference on Machine Learning*. PMLR, 2020, pp. 6316–6326.
- [33] A. Gretton, K. M. Borgwardt, M. J. Rasch, B. Schölkopf, and A. Smola, "A kernel two-sample test," *The Journal of Machine Learning Research*, vol. 13, no. 1, pp. 723–773, 2012.
- [34] J. You, R. Ying, X. Ren, W. Hamilton, and J. Leskovec, "Graphrnn: Generating realistic graphs with deep auto-regressive models," in *International Conference on Machine Learning*. PMLR, 2018, pp. 5708–5717.
- [35] E. K. Jackson, T. Levendovszky, and D. Balasubramanian, "Reasoning about metamodeling with formal specifications and automatic proofs," in *International Conference on Model Driven Engineering Languages and Systems*. Springer, 2011, pp. 653–667.
- [36] C. A. González, F. Büttner, R. Clarisó, and J. Cabot, "Emftocsp: A tool for the lightweight verification of emf models," in *2012 First International Workshop on Formal Methods in Software Engineering: Rigorous and Agile Approaches (FormSERA)*. IEEE, 2012, pp. 44–50.
- [37] N. Nassar, J. Kosiol, T. Kehrer, and G. Taentzer, "Generating large emf models efficiently," in *International Conference on Fundamental Approaches to Software Engineering*. Springer, 2020, pp. 224–244.
- [38] K. Ehrig, J. M. Küster, and G. Taentzer, "Generating instance models from meta models," *Software & Systems Modeling*, vol. 8, no. 4, pp. 479–500, 2009.
- [39] G. Soltana, M. Sabetzadeh, and L. C. Briand, "Synthetic data generation for statistical testing," in *2017 32nd IEEE/ACM International Conference on Automated Software Engineering (ASE)*. IEEE, 2017, pp. 872–882.
- [40] P. Erdős and A. Rényi, "On the evolution of random graphs," *Publ. Math. Inst. Hung. Acad. Sci.*, vol. 5, no. 1, pp. 17–60, 1960.
- [41] D. J. Watts and S. H. Strogatz, "Collective dynamics of 'small-world' networks," *nature*, vol. 393, no. 6684, pp. 440–442, 1998.
- [42] M. Simonovsky and N. Komodakis, "Graphvae: Towards generation of small graphs using variational autoencoders," in *International Conference on Artificial Neural Networks*. Springer, 2018, pp. 412–422.
- [43] Y. Li, O. Vinyals, C. Dyer, R. Pascanu, and P. Battaglia, "Learning deep generative models of graphs," *arXiv preprint arXiv:1803.03324*, 2018.